Systems biology

Simple sequence-based kernels do not predict protein–protein interactions

Jiantao Yu[1,2], Maozu Guo[1,*], Chris J. Needham[3], Yangchao Huang[1], Lu Cai[4] and David R. Westhead[2,*]

[1]School of Computer Science and Technology, Harbin Institute of Technology, Harbin 150001, China, [2]Institute of Molecular and Cellular Biology, [3]School of Computing, University of Leeds, Leeds, LS2 9JT, UK and [4]School of Mathematics, Physics and Biological Engineering, Inner Mongolia University of Science and Technology, Baotou, Inner Mongolia 014010, China

Associate Editor: Martin Bishop

ABSTRACT

Motivation: A number of methods have been reported that predict protein–protein interactions (PPIs) with high accuracy using only simple sequence-based features such as amino acid 3mer content. This is surprising, given that many protein interactions have high specificity that depends on detailed atomic recognition between physiochemically complementary surfaces. Are the reported high accuracies realistic?

Results: We find that the reported accuracies of the predictions are significantly over-estimated, and strongly dependent on the structure of the training and testing datasets used. The choice of which protein pairs are deemed as non-interactions in the training data has a variable impact on the accuracy estimates, and the accuracies can be artificially inflated by a bias towards dominant samples in the positive data which result from the presence of hub proteins in the protein interaction network. To address this bias, we propose a positive set-specific method to create a 'balanced' negative set maintaining the degree distribution for each protein, leading to the conclusion that simple sequence-based features contain insufficient information to be useful for predicting PPIs, but that protein domain-based features have some predictive value.

Availability: Our method, named '*BRS*-nonint', is available at http://www.bioinformatics.leeds.ac.uk/BRS-nonint/. All the datasets used in this study are derived from publicly available data, and are available at http://www.bioinformatics.leeds.ac.uk/BRS-nonint/PPI_RandomBalance.html

Contact: maozuguo@hit.edu.cn; d.r.westhead@leeds.ac.uk

1 INTRODUCTION

Protein–protein interactions (PPIs) are responsible for many critical functions and processes in biology, and are highly relevant to disease states. While experimental measurements are able to discover protein interactions at high throughput, they can be biased to specific types of interaction or unlikely to discover all interactions, and cannot be applied to all relevant species. As a result, a number of researchers have studied how well machine learning methods are able to predict PPIs from protein sequence information alone (Ben-Hur and Noble, 2005; Bock and Gough, 2001; Chou and Cai, 2006; Gomez *et al*., 2003; Guo *et al*., 2008, 2010; Martin *et al*., 2005; Nanni and Lumini, 2006; Park, 2009; Pitre *et al*., 2006, 2008; Roy *et al*., 2009; Shen *et al*., 2007; Sprinzak and Margalit, 2001; Yu *et al*., 2010). They have reported impressive performance using a variety of simple sequence features (such as counts of 3mers of neighbouring residues) and developed sophisticated pairwise kernels to allow classification using support vector machines (SVMs). The performance measures reported are much better than one would expect, and we have found that they are very sensitive to the content of the datasets used in training and cross-validation. In particular, the presence of 'hub' proteins that interact with many other proteins in the positive dataset leads to a strong bias that invalidates most performance estimates.

Recently, Park (2009) has benchmarked four methods on yeast and human data, found significant performance differences and highlighted some effects of training and testing data on performance estimates. However, the datasets used in this benchmark still suffer from bias of dataset content reported here. In this work, we present a method to remove this bias and demonstrate the effect on predicting PPIs from simple sequence features: that simple sequence-based kernels do not predict PPIs. We now summarize some approaches taken to predict PPIs from sequence data, all of which should be reinterpreted following our findings. Ben-Hur and Noble (2005) used a SVM method with a pairwise kernel and evaluated a number of sequence features, including the 'spectrum kernel' where the protein sequences are represented by the counts of 3mers of neighbouring residues (Leslie *et al*., 2002). This leads to a fixed length ($n=8000$) feature vector for each protein regardless of the sequence length. The pairwise kernel allows comparison of pairs of proteins, avoiding any problems that may be introduced by concatenating feature vectors for each protein, and on a yeast dataset performance of the SVM gives a receiver operator characteristic (ROC) area under curve (AUC) score of 0.81. Similarly Martin *et al*. (2005) used symmetric 3mers ($n=4200$) and a pairwise kernel giving accuracy of 70.3% on human data (and 69.0% for yeast). Accuracy of 83.9% was reported to be achieved by Shen *et al*. (2007) on human PPI data using 3mers of residues grouped into seven categories ($n=343$) and an S-kernel similar in function to the previous pairwise kernels. Guo *et al*. (2008, 2010)

---

*To whom correspondence should be addressed.

represented residues in seven categories with an auto-covariance method and reported high accuracies using shuffled sequences as a negative set, but poor results on randomly sampled negative data. Roy *et al.* (2009) proposed using simple amino acid composition features (normalized counts of single or pairs of amino acids) and reported results on a par with those obtained using protein domain information.

There exist a number of datasets of protein pairs which have been determined to interact by high-throughput methods (Costanzo *et al.*, 2010; Gavin *et al.*, 2006; Ito *et al.*, 2001; Stelzl *et al.*, 2005; Yu *et al.*, 2008) and can be used as positive data to train and validate prediction methods. In particular, many experiments have been conducted on yeast, and a PPI network consisting of 156 673 non-redundant interactions exists in BioGRID (Stark *et al.*, 2006), and a high-confidence dataset of 6568 non-redundant interactions has been published (Batada *et al.*, 2007). For human, the Human Protein References Database (HPRD) contains 38 945 non-redundant high-quality interactions (Peri *et al.*, 2004), and the BioGRID database contains 30 851 non-redundant interactions (Stark *et al.*, 2006). The choice of a set of negative examples (non-interacting protein pairs) to use in training and validating prediction methods is particularly important since few techniques can conduct a large-scale measurement of non-interacting pairs (Doerr, 2010; Smialowski *et al.*, 2010). The latest negative dataset, *Negatome* (Smialowski *et al.*, 2010), derives non-interactions from literature curation and protein complex structural data of the Protein Data Bank and can be seen as a meaningful start in the systematic construction of a set of reliable non-interactions in a biological context. However, each of the prediction methods reported above used a random sampling approach (sometimes modified by cellular localization information), selecting a negative set by random sampling of pairs of proteins that exist in the positive set but for which there is no evidence of interaction.

2 METHODS

2.1 Datasets

The datasets of interactions were processed to filter out redundant samples, self-interactions and proteins whose sequence contained elements other than the 20 standard amino acids. For yeast, 154 828 non-redundant positive interactions (47 335 physical interactions and remainder genetic) were extracted from the BioGRID database (Stark *et al.*, 2006) (3.0.64.mitab). For human, 36 134 non-redundant positive physical interactions were extracted from the HPRD (Peri *et al.*, 2004) (Release9_041310), which was used in Shen *et al.* (2007). In addition 27 307 human interactions extracted from the BioGRID database (3.0.64.mitab) were used to filter 'potential interactions' in creating negative candidates (see below) (16 993 PPIs overlap between these two sets). Sequence data used for yeast was from the *Saccharomyces Genome Database* (http://downloads.yeastgenome.org/) (file orf_trans.fasta), and for human was from UniProt (The UniProt Consortium, 2010) (file uniprot_sprot.fasta). For evaluating Pfam features, version 24.0 of the Pfam Database (Finn *et al.*, 2010) (pfam-A.hmm) was used. HMMER3.0 (http://hmmer.org) was used to scan the protein sequences. Smaller datasets were constructed to test the performance of Pfam features. For yeast, we used the high-confidence physical interaction dataset from Batada *et al.* (2007) (HC-BIOGRID-2.0.31.tab) and filtered out all interactions that involved a protein with no Pfam domain hits, resulting in a dataset containing 5621 interactions. For human, we took the intersection of the two human datasets described above, and again filtered out all interactions that involved a protein with no Pfam hits, resulting in a dataset containing 15 804 physical

interactions. We consider a 'Pfam hit' when the match has an *E*-value (full sequence) of less than 0.01, similar to the methods of Gomez *et al.* (2003).

2.2 Choosing a subset

Owing to the computational demand of the large datasets, we chose a subset of the interactions. Subsets were formed by choosing at random 300 proteins for yeast or 1500 for human, and then selecting all interactions involving any of those proteins. This creates an interaction dataset that maintains the vertex degree distribution of the large dataset. Typically, subsets contained ~15 000 interactions between ~4000 proteins for yeast and ~10 000 interactions between ~5300 proteins for human. We created 10 subsets at random and report average performances of 10-fold cross-validations (10×10 cv) in each of these subsets.

In order to test the stability of performance measures on subsets of different sizes, and also to examine the effect of excluding genetic interactions, we sampled four subsets of ~5000–~20 000 interactions from the yeast physical interactions in the BioGRID database (Stark *et al.*, 2006) (3.0.64.tab2). Ten cross-validations were performed on each subset, respectively.

2.3 Negative dataset construction

From the positive set of *N* interactions, the complement graph can be formed as the set of all possible pairs of proteins in the positive set for which a positive interaction is not present. Since we created subsets of the whole PPI network, we formed the negative candidate graph from the subset's complement graph. Interactions that exist in the BioGRID database (Stark *et al.*, 2006) were removed from the negative candidate graphs to reduce the potential for real interactions appearing erroneously in the negative sets. Self-interactions were also excluded. For *simple random sampling*, *N* edges were chosen from the negative candidate graph, ensuring that each protein that appeared in the positive dataset also appeared at least once in the negative dataset. Each protein was taken in turn and a connected edge randomly sampled from the negative candidate graph and added to the negative set until all *N* edges were chosen. For *balanced random sampling*, the number of times each protein appears in the negative set is equal to the number of times it appears in the positive set. Taking proteins in the positive dataset in turn, beginning with the protein of largest vertex degree, we randomly sampled edges from the negative candidate graph connected to the current protein until the vertex degree in the negative set was equal to that in the positive set.

The balanced random sampling method is available in the *BRS*-nonint software. Input to this software is the positive interaction dataset, as well as the pairs which should be avoided in choosing negative data. All the interactions (one pair per line) should be input with a format of 'proteinA proteinB interacting'.

2.4 SVMs and kernel functions

The methods of Ben-Hur and Noble (2005) and Shen *et al.* (2007) were evaluated, and the SVMs were implemented by modifying the code of libsvm-2.91 (http://www.csie.ntu.edu.tw/~cjlin/libsvm). From (Ben-Hur and Noble, 2005), the TPPK_{8000} kernel represents each protein as a length 8000 feature vector v_s of normalized counts (normalized to sum to one) of each possible triple of amino acids (the 3mer spectrum kernel). The kernel takes cosine form, $k_c(X, Y) = (v_{sx} \cdot v'_{sy})/sqrt((v_{sx} \cdot v'_{sx}) \times (v_{sy} \cdot v'_{sy}))$, for proteins *X* and *Y* with feature vectors v_{sx} and v_{sy}, respectively, and the pairwise kernel for proteins pairs (A–B and C–D) is calculated as a tensor product, $K_{\text{TPPK}}((A, B),(C, D)) = k_c(A,C) \times k_c(B, D) + k_c(A, D) \times k_c(B,C)$. From (Shen *et al.*, 2007), the S-kernel$_{343}$ represents each protein as a length 343 feature vector v_t of normalized counts (normalized by the largest) of each possible conjoint triad—3mers of residues grouped into seven categories. Squared Euclidean distances between proteins are calculated, $d_e(X, Y) = ||v_{tx} - v_{ty}||^2$ and the pairwise S-kernel is $K_S((A,B),(C,D)) = \exp(-\gamma||s||^2)$ with $s = \min(d_e(A, C) + d_e(B, D), d_e(A, D) + d_e(B, C))$. SVM parameters $C = 128$ and

$\gamma = 0.25$ were used as in Shen *et al.* (2007). For the Pfam kernel, a feature vector of counts of Pfam hits (see above) was created, and the TPPK pairwise kernel was used on these feature vectors (length 1932 and 3400 for yeast and human, respectively, as this many distinct Pfam features were detected). Similar results were obtained using other representations of the Pfam features (such as scores, *E*-values or binary presence/absence of domains).

2.5 Performance evaluation

Performances were reported by the AUC values (area under the ROC curve, a perfect classifier has AUC = 1, random performance gives AUC = 0.5) with 10-fold cross-validations. Variation in performance between 10 subsets or 10 cross-validations was quantified as 'average ± range'. Similar to ROC_{50} in Ben-Hur and Noble (2005), we also report $ROC_{(0.5\%)}$ (the area under the ROC curve for false positive rate $\leq 0.5\%$), which measures just the true positives detected before 0.5% of false positives.

3 RESULTS

The structures of the various datasets are illustrated in Figure 1, which clearly shows the hub proteins in the positive set, reflecting the well-known power-law distribution of vertex degrees, and how this is replicated in the negative set generated by the method of balanced random sampling but not simple random sampling.

3.1 Prediction performance and dataset bias

SVM performance results are reported in Table 1. The $TPPK_{8000}$ kernel (Ben-Hur and Noble, 2005) predicted PPIs with an AUC of 0.95 on the yeast dataset, when trained and tested with a negative set constructed by simple random sampling. We also reproduced the reported results of 0.81 AUC on the set of 10 517 yeast interactions used in Ben-Hur and Noble (2005). However, when a balanced negative dataset was used, the performance fell to 0.5 AUC, indicating that performance was close to random (AUC = 0.46 on the same dataset of Ben-Hur and Noble using balanced sampling). This highlights the bias inherent in the structure of the datasets constructed using a random sampling approach that does not account for the number of times each protein appears in the positive and negative datasets. This is also confirmed by similar results for predicting PPIs on human datasets (0.83 AUC versus 0.55 AUC on balanced dataset), again indicating the bias, and revealing the poor performance of the classifier based on these simple sequence features.

Using the S-kernel$_{343}$, we obtained performance of 0.67 AUC on human data with a negative set generated by simple random sampling and our own implementation of the methods of Shen *et al.* (2007). Using a negative set constructed by balanced sampling this was reduced to 0.55 AUC. On yeast data, the AUC score dropped from 0.72 to 0.53 when the balanced negative set was used.

The data in Table 2 show the effect of using different subset sizes, and restricting the yeast dataset to physical interactions only (NB: the human data in Table 1 already consist entirely of physical interactions). Similar to the above results, the AUC values drop from ~0.9 (simple random sampling) to 0.6 (balanced sampling), and these values are stable for different subset sizes. Comparing these values to the Table 1 results for yeast, shows that omitting yeast genetic interactions leads to AUC values about 0.05 lower for simple random sampling and 0.1 higher for balanced sampling. The decrease in AUC from simple random sampling when genetic interactions are omitted is probably caused by a decrease in dataset
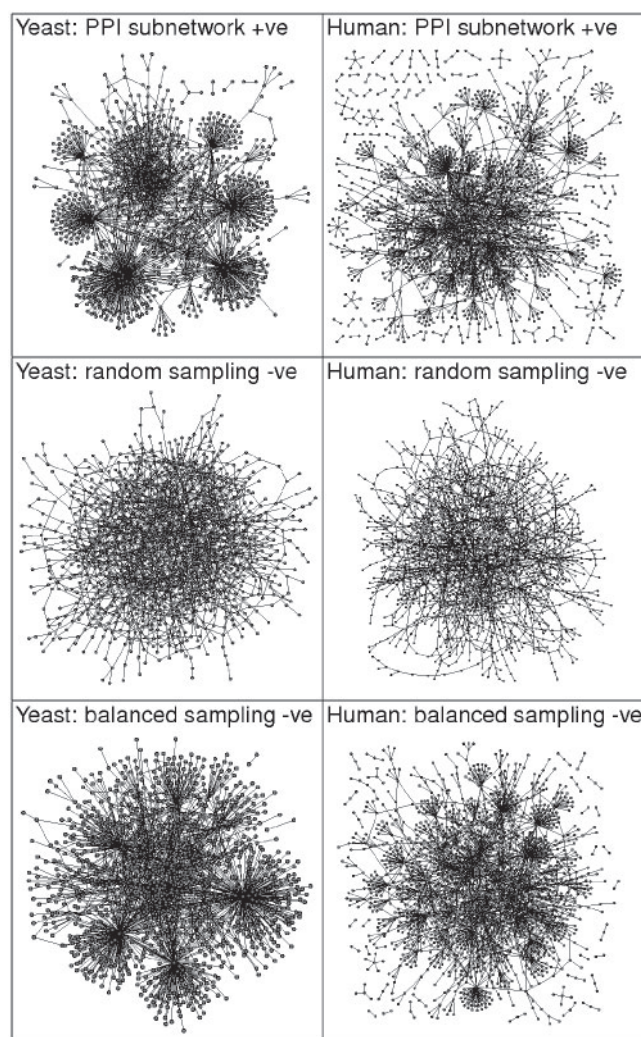

Fig. 1. PPI networks—visualization of the network structures, plotted by Cytoscape (Shannon *et al.*, 2003). A positive dataset, and negative datasets generated by *simple random sampling* and *balanced random sampling* method are presented on two subsets of the yeast and human PPI networks, generated as described in the text but containing only ~2000 interactions for visual clarity.

bias: genetic interactions lead to hubs of higher vertex degree in the positive dataset. The increase in AUC for balanced random sampling could perhaps indicate that physical interactions are slightly more predictable from simple sequence features than genetic interactions, but this is not borne out in all datasets (see for instance the smaller dataset of physical interactions used in Table 3, *vide infra*).

3.2 Predictions based on other sequence features

The above results suggest that the prediction of protein interactions from simple sequence features is very difficult. An alternative is to employ evolutionary information, in the form of recognizable conserved domains within the sequence, because it is known that some domains have propensities to interact with other domains, and their presence in the sequence could therefore be predictive of interactions to some extent. We used Pfam (Finn *et al.*, 2010) to

Table 1. Performance evaluation of classifier

Kernel		YEAST subsets ~15 000 interactions ~4000 proteins		HUMAN subsets ~10 000 interactions ~5300 proteins	
		Random sampling	Balanced sampling	Random sampling	Balanced sampling
TPPK$_{8000}$	ROC	0.95 ± 0.01	0.50 ± 0.03	0.83 ± 0.01	0.55 ± 0.01
	ROC$_{(0.5\%)}$	0.09	0.01	0.09	0.02
S-kernel$_{343}$	ROC	0.72 ± 0.02	0.53 ± 0.01	0.67 ± 0.02	0.55 ± 0.01
	ROC$_{(0.5\%)}$	0.05	0.00	0.03	0.00

Figures reported are the AUC values (area under the ROC curve) performed with 10-fold cross-validations on 10 subsets of the yeast and human datasets as detailed in Section 2. Values given are average ± range.

Table 2. Performance evaluation of classifiers

Kernel		YEAST subsets of different sizes			
		Number of interactions		Random sampling	Balanced sampling
TPPK$_{8000}$		~5000	ROC	0.92 ± 0.02	0.56 ± 0.02
			ROC$_{(0.5\%)}$	0.08	0.03
		~10 000	ROC	0.91 ± 0.01	0.59 ± 0.01
			ROC$_{(0.5\%)}$	0.07	0.03
		~15 000	ROC	0.88 ± 0.02	0.62 ± 0.01
			ROC$_{(0.5\%)}$	0.08	0.01
		~20 000	ROC	0.89 ± 0.01	0.63 ± 0.02
			ROC$_{(0.5\%)}$	0.09	0.02

Figures reported are the AUC values (area under the ROC curve) performed with 10-fold cross-validations on four subsets of yeast dataset which is composed of 47 335 physical interactions as detailed in Section 2. Values given are average ± range.

Table 3. Performance evaluation of classifiers

Kernel		YEAST HC PPI network 5621 interactions 2245 proteins		HUMAN HC PPI network 15 804 interactions 6198 proteins	
		Random sampling	Balanced sampling	Random sampling	Balanced sampling
Pfam$_{TPPK}$	ROC	0.74 ± 0.02	0.73 ± 0.02	0.77 ± 0.01	0.75 ± 0.01
	ROC$_{(0.5\%)}$	0.07	0.05	0.05	0.03
TPPK$_{8000}$	ROC	0.76 ± 0.02	0.48 ± 0.02	0.82 ± 0.01	0.60 ± 0.01
	ROC$_{(0.5\%)}$	0.03	0.01	0.07	0.02
S-kernel$_{343}$	ROC	0.64 ± 0.03	0.56 ± 0.03	0.68 ± 0.02	0.56 ± 0.01
	ROC$_{(0.5\%)}$	0.05	0.01	0.00	0.00

Figures reported are the AUC values (area under the ROC curve) performed with 10-fold cross-validations for the Pfam kernel and the spectrum kernels. Values given are average ± range.

identify domains and used a vector of counts of these domains in each sequence with the TPPK kernel to predict PPIs. We found that the Pfam domains are useful features for predicting PPIs, demonstrated by the result of 0.73 AUC on yeast and 0.75 AUC on human using a negative set generated by balanced sampling, and therefore contain stronger information about PPIs than we are able to extract from a spectrum kernel representation using current state of the art machine learning methods (see Table 3 for comparison with the TPPK$_{8000}$ and S-kernel$_{343}$ spectrum kernels on the same PPI network datasets). The Pfam features encode information that is useful in predicting PPIs, rather than encoding for particular proteins, demonstrated by a similar result of 0.74 AUC on yeast and 0.77 AUC on human using a negative set generated by simple random sampling: here there is minimal bias from the positive and negative dataset composition since the features allow the generalization of the important relationships between Pfam features present in protein sequences, rather than learning to identify that specific proteins tend to appear more in one set.

4 DISCUSSION

The high accuracy reported in the literature for predicting PPIs from simple sequence features appears to be an artefact of the datasets used to train and validate methods. In simple terms, if datasets are used where some proteins (hubs) appear many more times in the positive set than the negative set, then a machine learning method will learn this, predict positive interactions preferentially for these proteins and seem to be highly accurate. The sequence feature vectors used are of high dimension and allow machine learning methods to identify such specific proteins. The accuracies are unrealistic because the intended application is finding interactors for a specific protein from the entire proteome, of which even hub proteins only interact with a small fraction. It seems in fact that simple sequence features are little better than random in predicting protein interactions, when the task is to distinguish positive and negative interactions that occur in equal number for each protein in balanced positive and negative datasets.

It is clear that the choice of negative data is critical in training and performance evaluation for machine learning methods; in particular, there must be no systematic differences between positive and negative datasets used in training and evaluation that are not present in the datasets of intended application. While it is clear that the Negatome (Smialowski *et al.*, 2010) dataset will be useful in method development, it is also probable that this set will have its own biases, for instance towards well-studied proteins in the literature, and proteins that appear in the structural database. It will therefore have to be used with care in prediction method development.

While simple sequence features are not strongly predictive of interactions, the presence of conserved domains does contain some predictive information. But even then, accuracy is only moderate. Given that the a priori probability of interaction of a protein and a partner chosen randomly from the entire proteome is quite low, any of these methods will be susceptible to a high rate of false positive predictions. Protein interactions depend exquisitely on the three-dimensional atomic structures of the proteins concerned, and their detailed spatial and temporal pattern of expressions in the cell. Our ability to predict these interactions from information only in the sequence is, at present, too poor to be useful.

ACKNOWLEDGEMENTS

REFERENCES

Batada,N.N. *et al.* (2007) Still stratus not altocumulus: further evidence against the date/party hub distinction. *PLoS Biol.*, **5**, e154.

Ben-Hur,A. and Noble,W.S. (2005) Kernel methods for predicting protein-protein interactions. *Bioinformatics*, **21**, i38–i46.

Bock,J.R. and Gough,D.A. (2001) Predicting protein-protein interactions from primary structure. *Bioinformatics*, **17**, 455–460.

Chou,K.C. and Cai,Y.D. (2006) Predicting protein-protein interactions from sequences in a hybridization space. *J. Proteome Res.*, **5**, 316–322.

Costanzo,M. *et al.* (2010) The genetic landscape of a cell. *Science*, **327**, 425–431.

Doerr,A. (2010) The importance of being negative. *Nat. Methods*, **7**, 10–11.

Finn,R.D. *et al.* (2010) The Pfam protein families database. *Nucleic Acids Res.*, **38**, D211–D222.

Gavin,A.C. *et al.* (2006) Proteome survey reveals modularity of the yeast cell machinery. *Nature*, **440**, 631–636.

Gomez,S.M. *et al.* (2003) Learning to predict protein-protein interactions from protein sequences. *Bioinformatics*, **19**, 1875–1881.

Guo,Y. *et al.* (2008) Using support vector machine combined with auto covariance to predict protein-protein interactions from protein sequences. *Nucleic Acids Res.*, **36**, 3025–3030.

Guo,Y. *et al.* (2010) PRED_PPI: a server for predicting protein-protein interactions based on sequence data with probability assignment. *BMC Res. Notes*, **3**, 145.

Ito,T. *et al.* (2001) A comprehensive two-hybrid analysis to explore the yeast protein interactome. *Proc. Natl Acad. Sci. USA*, **98**, 4569–4574.

Leslie,C. *et al.* (2002) The spectrum kernel: a string kernel for SVM protein classification. *Proc. Pac. Symp. Biocomput.*, 564–575.

Martin,S. *et al.* (2005) Predicting protein-protein interactions using signature products. *Bioinformatics*, **21**, 218–226.

Nanni,L. and Lumini,A. (2006) An ensemble of K-local hyperplanes for predicting protein-protein interactions. *Bioinformatics*, **22**, 1207–1210.

Park,Y. (2009) Critical assessment of sequence-based protein-protein interaction prediction methods that do not require homologous protein sequences. *BMC Bioinformatics*, **10**, 419.

Peri,S. *et al.* (2004) Human protein reference database as a discovery resource for proteomics. *Nucleic Acids Res.*, **32**, D497–D501.

Pitre,S. *et al.* (2006) PIPE: a protein-protein interaction prediction engine based on the re-occurring short polypeptide sequences between known interacting protein pairs. *BMC Bioinformatics*, **7**, 365.

Pitre,S. *et al.* (2008) Global investigation of protein-protein interactions in yeast *Saccharomyces cerevisiae* using re-occuring short polypeptide sequences. *Nucleic Acids Res.*, **36**, 4286–4294.

Roy,S. *et al.* (2009) Exploiting amino acid composition for predicting protein-protein interactions. *PLoS ONE*, **4**, e7813.

Shannon,P. *et al.* (2003) Cytoscape: a software environment for integrated models of biomolecular interaction networks. *Genome Res.*, **13**, 2498–2504.

Shen,J. *et al.* (2007) Predicting protein-protein interactions based only on sequences information. *Proc. Natl Acad. Sci. USA*, **104**, 4337–4341.

Smialowski,P. *et al.* (2010) The negatome database: a reference set of non-interacting protein pairs. *Nucleic Acids Res.*, **38**, D540–D544.

Sprinzak,E. and Margalit,H. (2001) Correlated sequence-signatures as markers of protein-protein interaction. *J. Mol. Biol.*, **311**, 681–692.

Stark,C. *et al.* (2006) BioGRID: a general repository for interaction datasets. *Nucleic Acids Res.*, **34**, D535–D539.

Stelzl,U. *et al.* (2005) A human protein-protein interaction network: a resource for annotating the proteome. *Cell*, **122**, 957–968.

The UniProt Consortium. (2010) The Universal Protein Resource (UniProt) in 2010. *Nucleic Acids Res.*, **38**, D142–D148.

Yu,C.Y. *et al.* (2010) Predicting protein-protein interactions in unbalanced data using the primary structure of proteins. *BMC Bioinformatics*, **11**, 167.

Yu,H. *et al.* (2008) High-quality binary protein interaction map of the yeast interactome network. *Science*, **322**, 104–110.